

Data Augmentation

```
In [1]: from tensorflow.keras.preprocessing.image import ImageDataGenerator
        from tensorflow.keras.utils import array_to_img, img_to_array, load_img

        # Create an ImageDataGenerator instance
        datagen = ImageDataGenerator(
            rotation_range=40,
            width_shift_range=0.2,
            height_shift_range=0.2,
            shear_range=0.2,
            zoom_range=0.2,
            horizontal_flip=True,
            fill_mode='nearest'
        )
```

```
In [5]: img = load_img(r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg",target_
        img)
```

Out[5]:



```
In [6]: x = img_to_array(img) # this is a Numpy array with shape (3, 150, 150)
        x = x.reshape((1,) + x.shape) # this is a Numpy array with shape (1, 3, 150, 150)

        # the .flow() command below generates batches of randomly transformed images
        # and saves the results to the `preview/` directory
        i = 0
        for batch in datagen.flow(x, batch_size=1,
                                   save_to_dir=r"C:\Users\chitt\OneDrive\Desktop\Data Aug
                                   mentation\preview",
                                   save_to_dir_only=True):
            i += 1
            if i > 10:
                break # otherwise the generator would loop indefinitely
```

Prediction Of Images Using Keras Models

RESNET50

```
In [1]: from keras.utils import array_to_img, img_to_array, load_img
```

```
In [2]: img = load_img(r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg")  
img
```

Out[2]:



```
In [8]: import numpy as np  
from tensorflow.keras.applications.resnet50 import ResNet50, preprocess_input, d  
from tensorflow.keras.preprocessing import image  
  
# Load the ResNet-50 model pre-trained on ImageNet data  
model = ResNet50(weights='imagenet')  
  
# Load and preprocess the input image  
img_path = r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg" # Replace  
img = image.load_img(img_path, target_size=(224, 224))  
img_array = image.img_to_array(img)  
img_array = np.expand_dims(img_array, axis=0)  
img_array = preprocess_input(img_array)  
  
# Make predictions  
predictions = model.predict(img_array)  
  
# Decode and print the top-3 predicted classes  
decoded_predictions = decode_predictions(predictions, top=3)[0]  
print("Predictions:")  
for i, (imagenet_id, label, score) in enumerate(decoded_predictions):  
    print(f"{i + 1}: {label} ({score:.2f})") # Indented correctly  
  
# Optionally, you can obtain the class index for the top prediction  
top_class_index = np.argmax(predictions[0])  
print(f"\nTop Prediction Class Index: {top_class_index}")
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels.h5

102967424/102967424 ————— 21s 0us/step

1/1 ————— 4s 4s/step

Downloading data from https://storage.googleapis.com/download.tensorflow.org/data/imagenet_class_index.json

35363/35363 ————— 0s 1us/step

Predictions:

1: tiger (0.90)

2: tiger_cat (0.10)

3: zebra (0.00)

Top Prediction Class Index: 292

ResNet50V2

In [7]: `img = load_img(r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg", target_img`

Out[7]:



```
In [10]: import numpy as np
from tensorflow.keras.applications.resnet_v2 import ResNet50V2, preprocess_input,
from tensorflow.keras.preprocessing import image
# Load the ResNet-50 model pre-trained on ImageNet data
model = ResNet50(weights='imagenet')

# Load and preprocess the input image
img_path = r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg" # Replace
img = image.load_img(img_path, target_size=(224, 224))
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0)
img_array = preprocess_input(img_array)

# Make predictions
predictions = model.predict(img_array)

# Decode and print the top-3 predicted classes
decoded_predictions = decode_predictions(predictions, top=3)[0]
print("Predictions:")
for i, (imagenet_id, label, score) in enumerate(decoded_predictions):
    print(f"{i + 1}: {label} ({score:.2f})") # Indented correctly

# Optionally, you can obtain the class index for the top prediction
```

```
top_class_index = np.argmax(predictions[0])
print(f"\nTop Prediction Class Index: {top_class_index}")
```

1/1 ————— 4s 4s/step

Predictions:

1: nematode (0.10)
2: matchstick (0.04)
3: nail (0.02)

Top Prediction Class Index: 111

VGG16

In [8]: `img = load_img(r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg", target_img`

Out[8]:



```
In [13]: import numpy as np
from tensorflow.keras.applications.vgg16 import VGG16, preprocess_input, decode_
from tensorflow.keras.preprocessing import image

# Load the VGG16 model pre-trained on ImageNet data
model = VGG16(weights='imagenet')

# Load and preprocess the input image
img_path = r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg" # Replace
img = image.load_img(img_path, target_size=(224, 224))
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0)
img_array = preprocess_input(img_array)

# Make predictions
predictions = model.predict(img_array)

# Decode and print the top-3 predicted classes
decoded_predictions = decode_predictions(predictions, top=3)[0]
print("Predictions:")
for i, (imagenet_id, label, score) in enumerate(decoded_predictions):
    print(f"{i + 1}: {label} ({score:.2f})")

# Optionally, you can obtain the class index for the top prediction
top_class_index = np.argmax(predictions[0])
print(f"\nTop Prediction Class Index: {top_class_index}")
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels.h5

553467096/553467096 ————— 165s 0us/step

1/1 ————— 1s 988ms/step

Predictions:

1: tiger (0.78)

2: tiger_cat (0.22)

3: lynx (0.00)

Top Prediction Class Index: 292

VGG19

In [9]: `img = load_img(r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg",target_img`

Out[9]:



```
In [13]: import numpy as np
from tensorflow.keras.applications.vgg19 import VGG19, preprocess_input, decode_predictions
from tensorflow.keras.preprocessing import image

# Load the VGG19 model pre-trained on ImageNet data
model = VGG19(weights='imagenet')

# Load and preprocess the input image
img_path = r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg" # Replace
img = image.load_img(img_path, target_size=(224, 224))
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0)
img_array = preprocess_input(img_array)

# Make predictions
predictions = model.predict(img_array)

# Decode and print the top-3 predicted classes
decoded_predictions = decode_predictions(predictions, top=3)[0]
print("Predictions:")
for i, (imagenet_id, label, score) in enumerate(decoded_predictions):
    print(f"{i + 1}: {label} ({score:.2f})")

# Optionally, you can obtain the class index for the top prediction
top_class_index = np.argmax(predictions[0])
print(f"\nTop Prediction Class Index: {top_class_index}")
```


1/1 ————— 1s 1s/step

Predictions:

1: tiger (0.78)
2: tiger_cat (0.22)
3: zebra (0.00)

Top Prediction Class Index: 292

Xception

In [14]: `img = load_img(r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg",target_img`

Out[14]:



```
In [11]: import numpy as np
from tensorflow.keras.applications import Xception
from tensorflow.keras.applications.xception import preprocess_input, decode_predictions
from tensorflow.keras.preprocessing import image

# Load the Xception model pre-trained on ImageNet data
model = Xception(weights='imagenet')

# Load and preprocess the input image
img_path = r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg" # Replace
img = image.load_img(img_path, target_size=(299, 299)) # Xception expects 299x2
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0)
img_array = preprocess_input(img_array)

# Make predictions
predictions = model.predict(img_array)

# Decode and print the top-3 predicted classes
decoded_predictions = decode_predictions(predictions, top=3)[0]
print("Predictions:")
for i, (imagenet_id, label, score) in enumerate(decoded_predictions):
    print(f"{i + 1}: {label} ({score:.2f})")

# Optionally, you can obtain the class index for the top prediction
top_class_index = np.argmax(predictions[0])
print(f"\nTop Prediction Class Index: {top_class_index}")
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/xception/xception_weights_tf_dim_ordering_tf_kernels.h5

91884032/91884032 ————— 9s 0us/step

1/1 ————— 3s 3s/step

Predictions:

1: tiger (0.86)

2: tiger_cat (0.06)

3: jaguar (0.00)

Top Prediction Class Index: 292

InceptionV3

In [21]: `img = load_img(r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg",target_img`

Out[21]:



```
In [16]: import numpy as np
import time
from tensorflow.keras.applications import InceptionV3
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.inception_v3 import preprocess_input, decode_predictions

# Load the InceptionV3 model pre-trained on ImageNet data
model = InceptionV3(weights='imagenet')

# Load and preprocess the input image
img_path = r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg" # Replace
img = image.load_img(img_path, target_size=(299, 299)) # InceptionV3 expects 29
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0)
img_array = preprocess_input(img_array)

# Measure inference time
start_time = time.time()
predictions = model.predict(img_array)
end_time = time.time()

# Decode and print the top-5 predicted classes
decoded_predictions = decode_predictions(predictions, top=5)[0]
print("Predictions:")
for i, (imagenet_id, label, score) in enumerate(decoded_predictions):
    print(f"{i + 1}: {label} ({score:.2f})")
```

```

# Optionally, you can obtain the class index for the top prediction
top_class_index = np.argmax(predictions[0])
print(f"\nTop Prediction Class Index: {top_class_index}")

# Calculate and print the inference time per step
inference_time_ms = (end_time - start_time) * 1000.0
print(f"Inference Time: {inference_time_ms:.2f} ms")

# Model summary provides information about parameters and layers
# model.summary()

# Get the size of the model in megabytes
model_size_MB = model.count_params() * 4 / (1024 ** 2) # 4 bytes for float32
print(f"Size (MB): {model_size_MB:.2f} MB")

# Get the number of parameters and depth from the model's layers
num_parameters = model.count_params()
model_depth = len(model.layers)
print(f"Parameters: {num_parameters}")
print(f"Depth: {model_depth}")

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-application-s/inception_v3/inception_v3_weights_tf_dim_ordering_tf_kernels.h5

96112376/96112376 ————— 10s 0us/step

1/1 ————— 5s 5s/step

Predictions:

1: tiger (0.89)
 2: tiger_cat (0.04)
 3: lynx (0.00)
 4: zebra (0.00)
 5: jaguar (0.00)

Top Prediction Class Index: 292

Inference Time: 5045.19 ms

Size (MB): 90.99 MB

Parameters: 23851784

Depth: 313

MobileNetV2

```

In [22]: img = load_img(r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg",target_
img

```


Out[22]:



```
In [20]: import numpy as np
import time
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.mobilenet_v2 import preprocess_input, decode_

# Load the MobileNetV2 model pre-trained on ImageNet data
model = MobileNetV2(weights='imagenet')

# Load and preprocess the input image
img_path = r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg" # Replace
img = image.load_img(img_path, target_size=(224, 224)) # MobileNetV2 expects 22
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0)
img_array = preprocess_input(img_array)

# Measure inference time
start_time = time.time()
predictions = model.predict(img_array)
end_time = time.time()

# Decode and print the top-5 predicted classes
decoded_predictions = decode_predictions(predictions, top=5)[0]
print("Predictions:")
for i, (imagenet_id, label, score) in enumerate(decoded_predictions):
    print(f"{i + 1}: {label} ({score:.2f})")

# Optionally, you can obtain the class index for the top prediction
top_class_index = np.argmax(predictions[0])
print(f"\nTop Prediction Class Index: {top_class_index}")

# Calculate and print the inference time per step
inference_time_ms = (end_time - start_time) * 1000.0
print(f"Inference Time: {inference_time_ms:.2f} ms")

# Model summary provides information about parameters and layers
# model.summary()

# Get the size of the model in megabytes
model_size_MB = model.count_params() * 4 / (1024 ** 2) # 4 bytes for float32
print(f"Size (MB): {model_size_MB:.2f} MB")

# Get the number of parameters and depth from the model's layers
num_parameters = model.count_params()
```

```

model_depth = len(model.layers)
print(f"Parameters: {num_parameters}")
print(f"Depth: {model_depth}")

```

1/1 ————— 2s 2s/step

Predictions:

1: tiger (0.81)
 2: tiger_cat (0.06)
 3: jaguar (0.01)
 4: zebra (0.00)
 5: leopard (0.00)

Top Prediction Class Index: 292

Inference Time: 2125.15 ms

Size (MB): 13.50 MB

Parameters: 3538984

Depth: 156

MobileNetV2

In [24]: `img = load_img(r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg",target_img`

Out[24]:



```

In [23]: import numpy as np
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.mobilenet_v2 import preprocess_input, decode_p
# Load the MobileNetV2 model pre-trained on ImageNet data
model = MobileNetV2(weights='imagenet')

# Load and preprocess the input image
img_path = r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg" # Replace
img = image.load_img(img_path, target_size=(224, 224)) # MobileNetV2 expects 22
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0)
img_array = preprocess_input(img_array)

# Measure inference time
start_time = time.time()
predictions = model.predict(img_array)
end_time = time.time()

# Decode and print the top-5 predicted classes

```

```

decoded_predictions = decode_predictions(predictions, top=5)[0]
print("Predictions:")
for i, (imagenet_id, label, score) in enumerate(decoded_predictions):
    print(f"{i + 1}: {label} ({score:.2f})")

# Optionally, you can obtain the class index for the top prediction
top_class_index = np.argmax(predictions[0])
print(f"\nTop Prediction Class Index: {top_class_index}")
# Calculate and print the inference time per step
inference_time_ms = (end_time - start_time) * 1000.0
print(f"Inference Time: {inference_time_ms:.2f} ms")

# Model summary provides information about parameters and layers
# model.summary()

# Get the size of the model in megabytes
model_size_MB = model.count_params() * 4 / (1024 ** 2) # 4 bytes for float32
print(f"Size (MB): {model_size_MB:.2f} MB")

# Get the number of parameters and depth from the model's layers
num_parameters = model.count_params()
model_depth = len(model.layers)
print(f"Parameters: {num_parameters}")
print(f"Depth: {model_depth}")

```

WARNING:tensorflow:5 out of the last 5 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x0000028134164C10> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

1/1  2s 2s/step

Predictions:

1: tiger (0.81)
 2: tiger_cat (0.06)
 3: jaguar (0.01)
 4: zebra (0.00)
 5: leopard (0.00)

Top Prediction Class Index: 292


Inference Time: 1996.47 ms

Size (MB): 13.50 MB

Parameters: 3538984

Depth: 156

DenseNet121

In [26]: 

Out[26]:



```
In [25]: import numpy as np
from tensorflow.keras.applications import DenseNet121
from tensorflow.keras.applications.densenet import preprocess_input, decode_predictions
from tensorflow.keras.preprocessing import image
# Load the DenseNet121 model pre-trained on ImageNet data
model = DenseNet121(weights='imagenet')
# Load and preprocess the input image
img_path = r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg" # Replace
img = image.load_img(img_path, target_size=(224, 224)) # MobileNetV2 expects 22
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0)
img_array = preprocess_input(img_array)

# Measure inference time
start_time = time.time()
predictions = model.predict(img_array)
end_time = time.time()

# Decode and print the top-5 predicted classes
decoded_predictions = decode_predictions(predictions, top=5)[0]
print("Predictions:")
for i, (imagenet_id, label, score) in enumerate(decoded_predictions):
    print(f"{i + 1}: {label} ({score:.2f})")

# Optionally, you can obtain the class index for the top prediction
top_class_index = np.argmax(predictions[0])
print(f"\nTop Prediction Class Index: {top_class_index}")
# Calculate and print the inference time per step
inference_time_ms = (end_time - start_time) * 1000.0
print(f"Inference Time: {inference_time_ms:.2f} ms")
# Model summary provides information about parameters and layers
# model.summary()

# Get the size of the model in megabytes
model_size_MB = model.count_params() * 4 / (1024 ** 2) # 4 bytes for float32
print(f"Size (MB): {model_size_MB:.2f} MB")

# Get the number of parameters and depth from the model's layers
num_parameters = model.count_params()
model_depth = len(model.layers)
print(f"Parameters: {num_parameters}")
print(f"Depth: {model_depth}")
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/densenet/densenet121_weights_tf_dim_ordering_tf_kernels.h5

33188688/33188688 ————— 4s 0us/step

WARNING:tensorflow:6 out of the last 6 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x000002813648D000> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2), @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

1/1 ————— 8s 8s/step

Predictions:

1: tiger (0.97)
2: tiger_cat (0.02)
3: zebra (0.00)
4: jaguar (0.00)
5: leopard (0.00)

Top Prediction Class Index: 292

Inference Time: 8462.83 ms

Size (MB): 30.76 MB

Parameters: 8062504

Depth: 429

NASNetMobile

```
In [28]: img = load_img(r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg",target_
img
```

Out[28]:



```
In [27]: import numpy as np
from tensorflow.keras.applications import NASNetMobile
from tensorflow.keras.applications.nasnet import preprocess_input,decode_predict
from tensorflow.keras.preprocessing import image
# Load the DenseNet121 model pre-trained on ImageNet data
model = DenseNet121(weights='imagenet')
# Load and preprocess the input image
img_path = r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg" # Replace
img = image.load_img(img_path, target_size=(224, 224)) # MobileNetV2 expects 22
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0)
```



```

img_array = preprocess_input(img_array)

# Measure inference time
start_time = time.time()
predictions = model.predict(img_array)
end_time = time.time()

# Decode and print the top-5 predicted classes
decoded_predictions = decode_predictions(predictions, top=5)[0]
print("Predictions:")
for i, (imagenet_id, label, score) in enumerate(decoded_predictions):
    print(f"{i + 1}: {label} ({score:.2f})")

# Optionally, you can obtain the class index for the top prediction
top_class_index = np.argmax(predictions[0])
print(f"\nTop Prediction Class Index: {top_class_index}")
# Calculate and print the inference time per step
inference_time_ms = (end_time - start_time) * 1000.0
print(f"Inference Time: {inference_time_ms:.2f} ms")
# Model summary provides information about parameters and layers
# model.summary()

# Get the size of the model in megabytes
model_size_MB = model.count_params() * 4 / (1024 ** 2) # 4 bytes for float32
print(f"Size (MB): {model_size_MB:.2f} MB")

# Get the number of parameters and depth from the model's layers
num_parameters = model.count_params()
model_depth = len(model.layers)
print(f"Parameters: {num_parameters}")
print(f"Depth: {model_depth}")

```

1/1  9s 9s/step

Predictions:

```

1: tiger (0.95)
2: tiger_cat (0.04)
3: zebra (0.00)
4: jaguar (0.00)
5: lynx (0.00)

```

Top Prediction Class Index: 292


Inference Time: 8761.78 ms

Size (MB): 30.76 MB

Parameters: 8062504

Depth: 429

NASNetLarge

In [30]:  `img = load_img(r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg", target_`
`img`

Out[30]:



```
In [29]: import numpy as np
from tensorflow.keras.applications import NASNetLarge
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.nasnet import preprocess_input, decode_predictions
# Load the DenseNet121 model pre-trained on ImageNet data
model = DenseNet121(weights='imagenet')
# Load and preprocess the input image
img_path = r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg" # Replace
img = image.load_img(img_path, target_size=(224, 224)) # MobileNetV2 expects 22
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0)
img_array = preprocess_input(img_array)

# Measure inference time
start_time = time.time()
predictions = model.predict(img_array)
end_time = time.time()

# Decode and print the top-5 predicted classes
decoded_predictions = decode_predictions(predictions, top=5)[0]
print("Predictions:")
for i, (imagenet_id, label, score) in enumerate(decoded_predictions):
    print(f"{i + 1}: {label} ({score:.2f})")

# Optionally, you can obtain the class index for the top prediction
top_class_index = np.argmax(predictions[0])
print(f"\nTop Prediction Class Index: {top_class_index}")
# Calculate and print the inference time per step
inference_time_ms = (end_time - start_time) * 1000.0
print(f"Inference Time: {inference_time_ms:.2f} ms")
# Model summary provides information about parameters and layers
# model.summary()

# Get the size of the model in megabytes
model_size_MB = model.count_params() * 4 / (1024 ** 2) # 4 bytes for float32
print(f"Size (MB): {model_size_MB:.2f} MB")

# Get the number of parameters and depth from the model's layers
num_parameters = model.count_params()
model_depth = len(model.layers)
print(f"Parameters: {num_parameters}")
print(f"Depth: {model_depth}")
```

1/1 ————— 10s 10s/step

Predictions:

1: tiger (0.95)
 2: tiger_cat (0.04)
 3: zebra (0.00)
 4: jaguar (0.00)
 5: lynx (0.00)

Top Prediction Class Index: 292

Inference Time: 10102.48 ms

Size (MB): 30.76 MB

Parameters: 8062504

Depth: 429

EfficientNetV2B0

In [31]: `img = load_img(r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg", target_img`

Out[31]:



```
In [32]: import numpy as np
from tensorflow.keras.applications import EfficientNetV2B0
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.efficientnet_v2 import preprocess_input, decode
# Load the DenseNet121 model pre-trained on ImageNet data
model = DenseNet121(weights='imagenet')
# Load and preprocess the input image
img_path = r"C:\Users\chitt\OneDrive\Pictures\image\Tiger Image.jpg" # Replace
img = image.load_img(img_path, target_size=(224, 224)) # MobileNetV2 expects 22
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0)
img_array = preprocess_input(img_array)

# Measure inference time
start_time = time.time()
predictions = model.predict(img_array)
end_time = time.time()

# Decode and print the top-5 predicted classes
decoded_predictions = decode_predictions(predictions, top=5)[0]
print("Predictions:")
for i, (imagenet_id, label, score) in enumerate(decoded_predictions):
    print(f"{i + 1}: {label} ({score:.2f})")
```

```
# Optionally, you can obtain the class index for the top prediction
top_class_index = np.argmax(predictions[0])
print(f"\nTop Prediction Class Index: {top_class_index}")
# Calculate and print the inference time per step
inference_time_ms = (end_time - start_time) * 1000.0
print(f"Inference Time: {inference_time_ms:.2f} ms")
# Model summary provides information about parameters and layers
# model.summary()

# Get the size of the model in megabytes
model_size_MB = model.count_params() * 4 / (1024 ** 2) # 4 bytes for float32
print(f"Size (MB): {model_size_MB:.2f} MB")

# Get the number of parameters and depth from the model's layers
num_parameters = model.count_params()
model_depth = len(model.layers)
print(f"Parameters: {num_parameters}")
print(f"Depth: {model_depth}")
```

1/1  9s 9s/step

Predictions:

1: matchstick (0.93)
2: honeycomb (0.06)
3: panpipe (0.01)
4: maraca (0.00)
5: web_site (0.00)

Top Prediction Class Index: 644

Inference Time: 9346.20 ms

Size (MB): 30.76 MB

Parameters: 8062504

Depth: 429

Completed

In []: